

Classes 1 and 2, January 25 and 30, 2024

We begin by describing the purely functional sublanguage of R5RS-Scheme

'Purely functional' means no ASSIGNMENT and, more generally, no 'destructive programming': after initial definition, values are not changed (or 'destroyed').

We focus on what is called DECLARATIVE PROGRAMMING in contrast to the IMPERATIVE PROGRAMMING style of

At the basis of this language are a few <sup>prev. class</sup> primitive data types. Initially, we'll make use of just three of these:

- numbers
- booleans
- functions

Numbers can be integers, rationals, or reals (complex, too, but we won't use these)

integers: 'infinite' precision

rationals:  $1/3$  displays as  $1/3$ , not as  $0.33333\ldots$

reals: subject to same precision limitations as in C

Booleans are just `#t` (true) and `#f` (false) — not themselves defined

This is one difference between R5RS and the version of MIT Scheme used in Abelson & Sussman (SICP). They use `nil` for `#f`, as well as for the empty list. Note that `nil` is not even defined in R5RS.

Functions come in two main flavors — primitives and user-defined. For this introductory look at the basis of R5RS, we'll focus on the arithmetic primitives.

Examples of these are `+`, `*`, `/`, `-`, `modulo`, `quotient`, `remainder`, ...

I urge you to get in the habit of taking incremental reads of the R5RS manual, where you will find a complete list of the primitives.

Numbers, Booleans and primitive functions can be said to be self-evaluating: type one in, and get it back, unchanged.

We can (of course) use these primitives to form more complex expressions

Before showing some of these, let me emphasize the critical role of parentheses in Scheme. Unlike their usage in math, parens are NOT just for the convenience of the reader.

When we write

$$(+ \ 2 \ 3)$$

for example, the parens are regarded as an instruction:

- evaluate the leftmost entry - throw an error if its value is not a function
- evaluate the remaining entries
- apply the value of the leftmost entry to these values
- return the result

For  $(+ 2 3)$ ,  $+$  evals to the built-in  $+$  function,  $2$  evals to the internal representation of  $2$ , and  $3$  to the internal rep. of  $3$ . The application then returns  $5$ .

Similarly for  $*$ ,  $-$ ,  $/$ ,  $\dots$

$(- 5 2)$  returns  $3$

$(/ 10 2)$  returns  $5$

and so on.

Interestingly, these primitives are generic in at least two ways.

First (unlike C),  $+$  is happy to accept reals and rationals as well as integers; indeed, even mixtures of these

(☐ will record that the lecture moved over to racket for some examples.)

- main takeaways from first excursion  
→ generic as regards numeric types  
(i.e., we do not need separate functions  
 $+int$ ,  $+real$ ,  $+rat$  ... )  
and allows mixed types  
→ despite the absence of  
user-created type tags, types  
are nonetheless present: one  
cannot compute  $(+ \#t \#f)$ , for  
example — the primitive  $+$   
requires that its arguments be  
numbers.

dynamic  
typing  
— types detected  
at run time,  
rather than at  
compile time

A second kind of genericity has to do with  
the number of arguments.

- main takeaways from second excursion  
→ generic also as regards the number  
of arguments. Eg  $(+ 1 2 3)$   
computes 6. The idea is that  $+$   
regards its args as a list, and it returns  
the sum of the numbers in that list.  
On this model,  $(+)$  ought also

make sense, and indeed  $(+)$  returns 0 — the sum of all the numbers in the empty list.

Similarly for the other arithmetic primitives? Explore!

Note that  $(*)$  returns 1, not 0, in line with mathematical convention.

Even more complex expressions can be formed if we nest parens. For example,

$(+ (* 3 4) 2)$

is eval'd using the same rule


(The Scheme Evaluation Rule)

- $+$  evals to the built-in  $+$  function (call it plus)
- $2$  evals to two


- The rule is applied recursively to get the value 12 for  $(\times 3 4)$ .

$$\text{So } (+ (\times 3 4) 2) = 14$$

There is no limit on the nesting depth, but as these expressions become larger, it will be helpful to have a simple visualization for them. We use trees in the expected way:-



It is important to know that one should not expect that  $+$  evaluates its args in left to right order. In fact, we know nothing at all about the eval. order



except that the arguments to a function (as well as the function itself) must have been evaluated before the function can be applied.

In fact, this is the ONLY notion of order which matters for this kind of programming. As you will soon see, this explains why the semicolon `(;)` has in Scheme



(or any functional language) nothing like the importance it has in C. Essentially, 'before and after' are concerns in Scheme only as they apply to the evaluation of a function and its args before the function can be applied to these arguments.

In other words, " $P; Q$ " (do  $P$ , then  $Q$ ) makes NO sense in (pure functional) Scheme !!!

*constructed as*

*We'll have a lot more to say about this.*

We will see how function composition makes up for this: assuming

$f$  and  $g$  are composable functions,  
 $(f \circ g)(x) = f(g(x))$

Next, let's look at  
variables.

We'll get at variables via the special form define.

For example

(define x 3)

We say that define is a special form because the 'Scheme Evaluation Rule' discussed last time does not apply for define.


So - how does define work?

Initially, scheme opens a global namespace in which the primitives (etc) are defined, but no user variables are defined.




initial (empty)  
environment

→ (define x 3) →



We write  $x:3$   
to indicate that  $x$   
is a name of the value 3


We say "x is bound to 3". Another suggestive picture is




If we were now to

(define y 3)

Then the global env would be



with




as the corresponding pointer diagram.

Once (define x 3) has been evaluated, then, within the scope of this define, you may replace all occurrences of x by 3 without changing the meaning of the substituted expression

(define x 3)

~ Scope of this definition



means (+ 3 1) : anywhere  $x$  occurs, we can replace it with a 3. This illustrates REFERENTIAL TRANSPARENCY

So what? Well, consider the contrast with C. After the declaration

$\text{int } x = 3$


consider

$x = x + 1.$


Can we replace  $x$  by 3 without wrecking the meaning? Clearly NOT!

What's responsible for the difference?

In C,  $x$  is NOT a name for 3 after this initialization — rather it is a name for a box (memory location) containing 3.




in C





in Scheme

(\*) In scheme, as in C, scopes frequently have "holes", typically brought about by local variables




One way of changing  $x$  and its scope:


(define  $x$  3)



(define  $x$  4)



PREVIOUS  
pointer  
disappears



Hey! Why isn't this just assignment?

Well, it is: it is definitional assignment, making an association between  $x$  and its value.

But it is extremely restricted, and Scheme in no way allows anything

equivalent to  $x = x + 1$ . In particular, `define` cannot occur within an expression — nor in the body of a program.

What haven't we explained yet about `define`?

What about

`(define x (+ x 1))`

?

This would assume that  $x$  has already been defined: The eval rule for

`(define x e)`

some perhaps complicated expression

is something like

- ① install  $x$  in the global env, perhaps without a value



$x: ?$

② Evaluate  $e$  to get its value, say  $v$

③ Finish the binding — so that, when the eval. of the def is complete, we would have

$x: v$  , i.e.,  $x \rightarrow v$

(Take a moment to think why the standard evaluation rule would not work for define.)

(define x e)

The standard rule would say

standard Rule

breaks  
here! →

- eval. the primitive define

- eval x (but x is usually  
here being defined  
for the first time)

so x has  
no value!

- eval e

Once a variable has been bound to a value, we can use it anywhere we like as (eg) a shorthand for that value: after

(define pi 3.1415 ---)

we could write (+ 2 pi)

(+ 2 (\* 3 pi))  
.  
.  
.

---

Another example

(define x 17)

could we then write

(define y x) ?

YES!

y: 17  
x: 17

---

- ① We talked about the special form `define` - but although I mentioned the severe restrictions on the use of `define`, I still need to give you an example of forbidden usage of `define` in an expression

Here's one

```
(+ (define x 3) 2)
```

You can use the Check Syntax button in dracket to see the explanation of why this won't work.

---

Recall the C distinction between expressions and statements. Pure functional scheme does not have statements.

So this restriction on `define` is very strong!

---

Should also have mentioned that `'define'` may not be used as a variable - it is a reserved word (\*\*) )

(\*\*) Just about any combination of numbers and characters can be used to form identifiers in Scheme. There are some restrictions

→ no embedded white space

→ no embedded parens

→ perhaps others, but I am not ~~\*~~ thinking of them just now

By default, scheme does not differentiate upper and lower case letters.

\* You will find the complete set of rules in the R5RS reference manual, which I have posted on Teams.

Another special form is

lambda

When we enter

(lambda (x) (\* x x))

in R5RS, the system returns

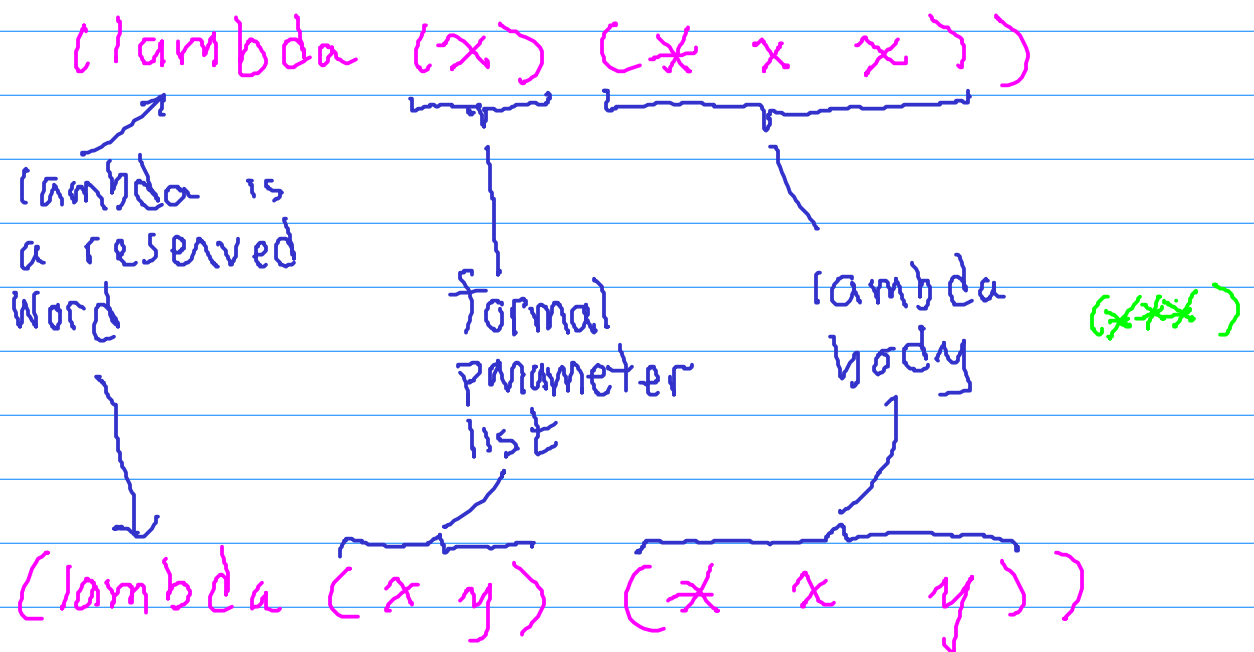
#procedure

We conclude that the lambda form is a 'stand-alone' scheme value.

In fact, such forms are scheme values in the same sense that, eg, numbers and booleans are scheme values. We will see that, just as numbers and booleans can be returned by function calls, so too can functions be returned by function calls. In fact, functions can occur as components of data structures — eg — we will see lists of functions.

One says that, in scheme, functions are first class.

It's useful to take the lambda form apart:



The formal parameters each open a <sup>local</sup> scope, which is precisely the body of the lambda form.

Returning to the Lecture 1 discussion, these scopes can be said to be holes in any containing scope of the same variable.

For example:

scope of the  $x$

```
(define x 1)
```

$x$  — This  $x$  is 1

scope of the formal parameter  $x$

```
((lambda (x) (+ x 2))) 17
```

$x$  — This  $x$  is 1      This  $x$  is 17

scope of  $x$  now with value 6

```
(define x 6)
```

You will want to experiment .. consider

```
(define x 1)
```

substitution rule

```
((lambda (x) (+ x 2))) (+ x 3)
```

$x$

what is the value of this  $x$ ?

The value is 1

value returned by this application

$$is (+ 4 2) = 6$$



( Anything of the form

(fn arg1 arg2 ... argk)  
is said to be an application. )

Note that define and lambda play nicely together :

(define square  
 (lambda (x) (\* x x)))

Abelson & Sussman notation for this function definition is

(define (square x)  
 (\* x x))

They prefer to not write the lambda.  
(This is the first instance we've seen  
of what is called syntactic sugar )

Having defined square this way, one uses it as follows

(square 2)

By the scheme eval-rule, square would be evaluated to

(lambda (x) (\* x x))

and then the evaluation continues as we have described, via the

substitution rule [substitute the value of the actual parameter, 2, for the formal parameter, x, everywhere x occurs in the lambda body.]

Still to discuss in connection with functions :

- type checking
- what is a function value?
- can we have an operational description of function evaluation?
- what happens if there is an undefined variable in a lambda body?
- how can functions be composed?
- your questions

Mentioned in class while using dracket :

- newline
- display

These primitives can be useful in understanding scheme - allowing simple instrumentation of our programs (eventually) and (now) of our interactions with the language.

I will summarize my use of them momentarily. first, however, it is important to know that neither of the calls

(newline)

(display e)

(where I assume  $e$  has been given a value previously) returns a value which can be input to another function.

If one attempts

`(+ (display 1) 2)`

for example an error will be thrown - complaining about the fact that `(display 1)` returns no useable value. All it does is print 1 to the screen (that's the side-effect) (\*)

That being said, here is an annotated copy of an interaction with racket

`(define x 1)` ; entered into a fresh window

`(newline)`  
`(display x)` ; see what happens when  
`(newline)` ; you remove the (newline)s.

x is 2 in the local scope

`((lambda (x y) (* (+ x (* y 2)) x)) 2 4)` ; opens a local scope for x and y

`(newline)`  
`(display x)` ← ; local scope has disappeared  
`(newline)` ; and x is again 1

`(define x 17)`

`(newline)` ; now x is 17 in the global  
`(display x)` ; scope: there is no record  
`(newline)` ; of it ever having been 1

Similarly, there is no record of x ever having been 2 ; you cannot "reach back into a scope"

What if, at the end of this session, we were to ask for a display of the value of  $y$ ? An error message, complaining about  $y$  being undefined, would be shown.

What about writing just  $y$ , instead of  $(\text{display } y)$ ? Same error: the first thing the special form `display` does is to evaluate its argument.

Note that `display` takes only a single argument. R5RS does not have a primitive corresponding to `printf`.

(\*) You are familiar with side-effects in C, though it is possible that this term was not used to describe them.

Consider what you can do with a function parameter that has been passed by reference: from inside the function, you can change the ("external") value of that parameter.

Such a change is said to be a side-effect of the function call.

(~~\*)~~) There is no requirement that all of the formal parameters actually occur in the lambda body. Eg

$((\text{lambda}(x\ y)\ x)\ 17\ 18)$

makes perfect sense, and returns 17.

Indeed, there is no requirement that ANY of the formal parameters occur in the function body! Eg

$(\text{lambda } (x) 3) 8$  will return 3.

This is said to be a constant function.

perhaps a better way to write such a function might be

$(\text{lambda } () 3)$

with calls having 0 actual parameters

$((\text{lambda } () 3))$

Note also that we may not write lambda forms with duplicated parameters — e.g.

$(\text{lambda } (x x) x)$

is not accepted by the reader.

(~~\*\*\*~~) I will use standard terminology to describe parameters: in the call

$(\text{lambda } (x \ y) (+ \ x \ y)) \ 2 \ 3)$

$x$  and  $y$  are the formal parameters and  $2$  and  $3$  are the actual parameters

(~~\*\*\*\*~~) The number of actual parameters is set by the definition of the function, and any call with a different number of actual params than the number of formal params will be flagged as an error. This is true for C++, as well: it is a type error.